

The Outliers

Who are we?

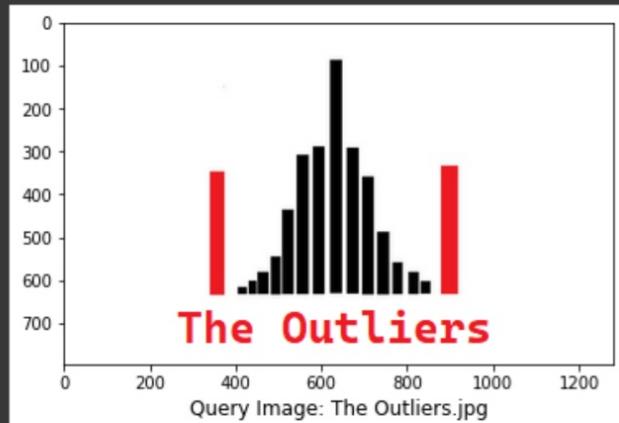
Why are we here?

The team

Problem Statement

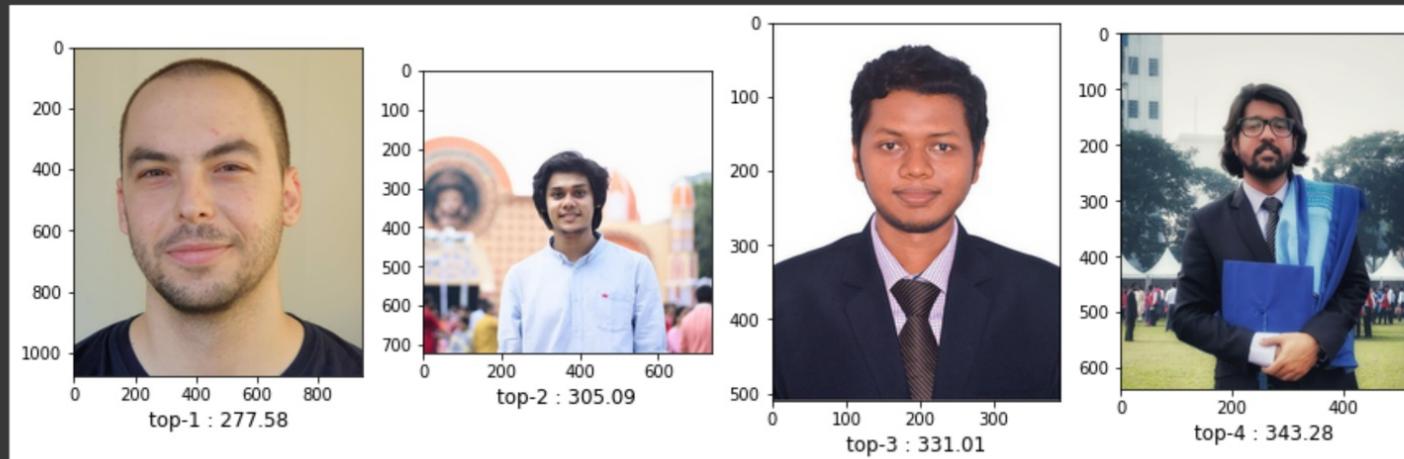
Related Works

----- Query Image -----



The Team

----- Similar Images -----



The Outliers

Who are we?

Why are we here?

The
team

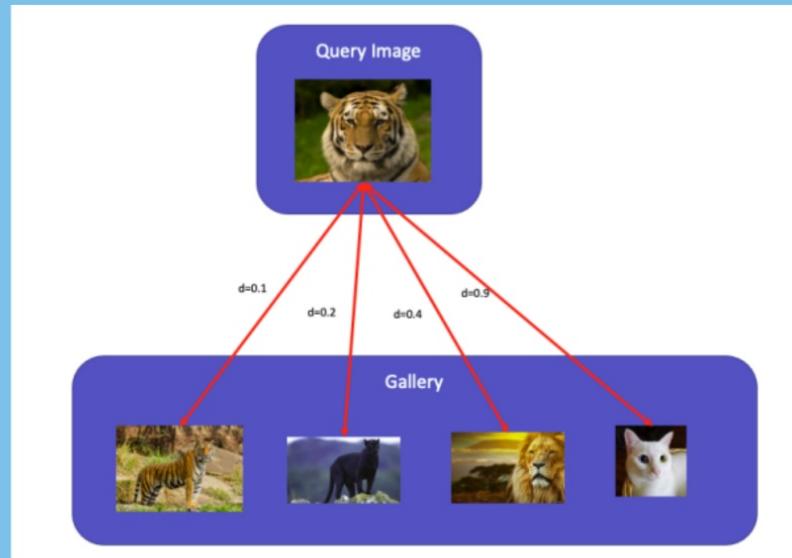
Problem
Statement

Related
Works

Problem Statement

*Create an image search engine where a **query image** is fed to a model that will return the most N similar images from a **gallery**.*

Given the input query image, the algorithm has to be capable of matching the input query image with another gallery image depicting the same animal.



The Outliers

Who are we?

Why are we here?

The team

Problem Statement

Related Works

"Good artists copy, great artists steal" Pablo Picasso

Ramprasath, Muthukrishnan, M. Vijay Anand, and Shanmugasundaram Hariharan.
"Image classification using convolutional neural networks." International Journal of Pure and Applied Mathematics 119.17 (2018): 1307-1319.

Chauhan, Rahul, Kamal Kumar Ghanshala, and R. C. Joshi. "**Convolutional neural network (CNN) for image detection and recognition.**" 2018 First International Conference on Secure Cyber Computing and Communication (ICSCCC). IEEE, 2018.

Hijazi, Samer, Rishi Kumar, and Chris Rowen.
"Using convolutional neural networks for image recognition." Cadence Design Systems Inc.: San Jose, CA, USA 9 (2015).

The Outliers

Who are we?

Why are we here?

The
team

Problem
Statement

Related
Works

Introduction to Machine Learning

"The Outliers"

1.
Introduction

2.
Theoretical
Description

3.
Data Collection

4.
Models
&
Implementations

5.
Results
&
Challenges

6.
Conclusion

Theories & Techniques

Approach
to the
Problem

Models

Motivations

Image Recognition & Classification

- Make some **research** and pick the most popular **solutions** & later **compare them**;
- Get and **pre-process data**;
- **Implement** and try all the possible solutions;
- Fit our work to make it feasible for the **competition**.



Theories & Techniques

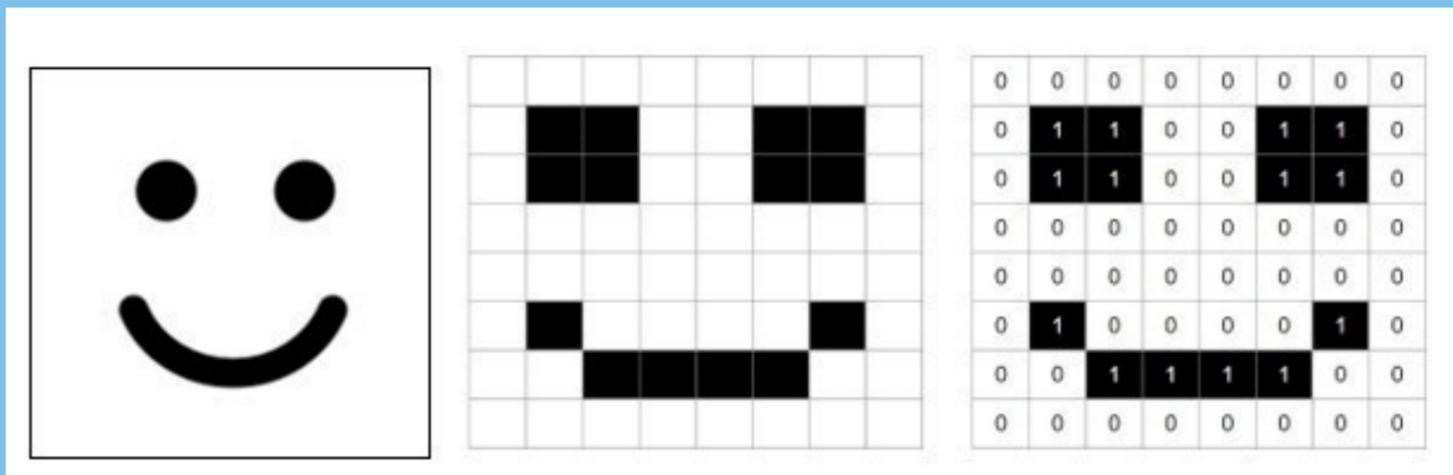
Approach
to the
Problem

Models

Motivations

CNN - Convolutional Neural Network

A Convolutional Neural Network, also known as CNN, is a class of **neural networks** that **specializes** in **processing data** that has a **grid-like topology**, such as an image.



Theories & Techniques

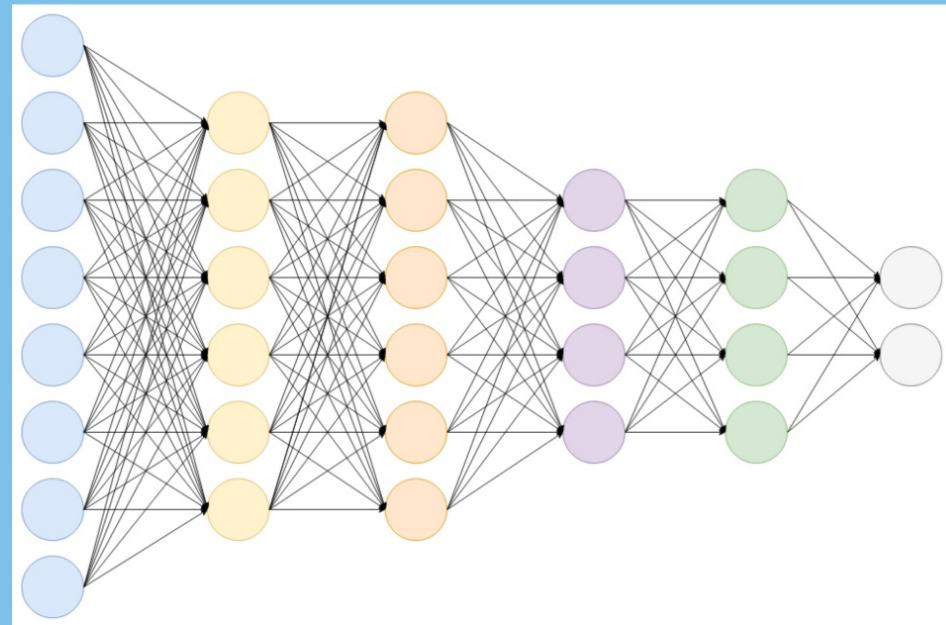
Approach
to the
Problem

Models

Motivations

Why did we choose it among all?

- 1- CNNs is **specialised** in **pattern detection**;
- 2- There is a lot of **documentation** about CNNs;
- 3- There are many **pretrained models** which are applicable to this project problem.



Theories & Techniques

Approach
to the
Problem

Models

Motivations

Introduction to Machine Learning

"The Outliers"

1.
Introduction

2.
Theoretical
Description

3.
Data Collection

4.
Models
&
Implementations

5.
Results
&
Challenges

6.
Conclusion



Here comes the miners:
Data Collection

Image Scraping

Amount of Data

Characteristics
of the Data

Our Sources

- Image Download Tool
- Web Scraping
- Data Cleaning
- <https://www.kaggle.com/datasets/iamsouravbanerjee/animal-image-dataset-90-different-animals/datasets>
- <https://www.image-net.org/challenges/LSVRC/2014/index#data>
- <https://www.mediasearch.dev/image>

Approaches

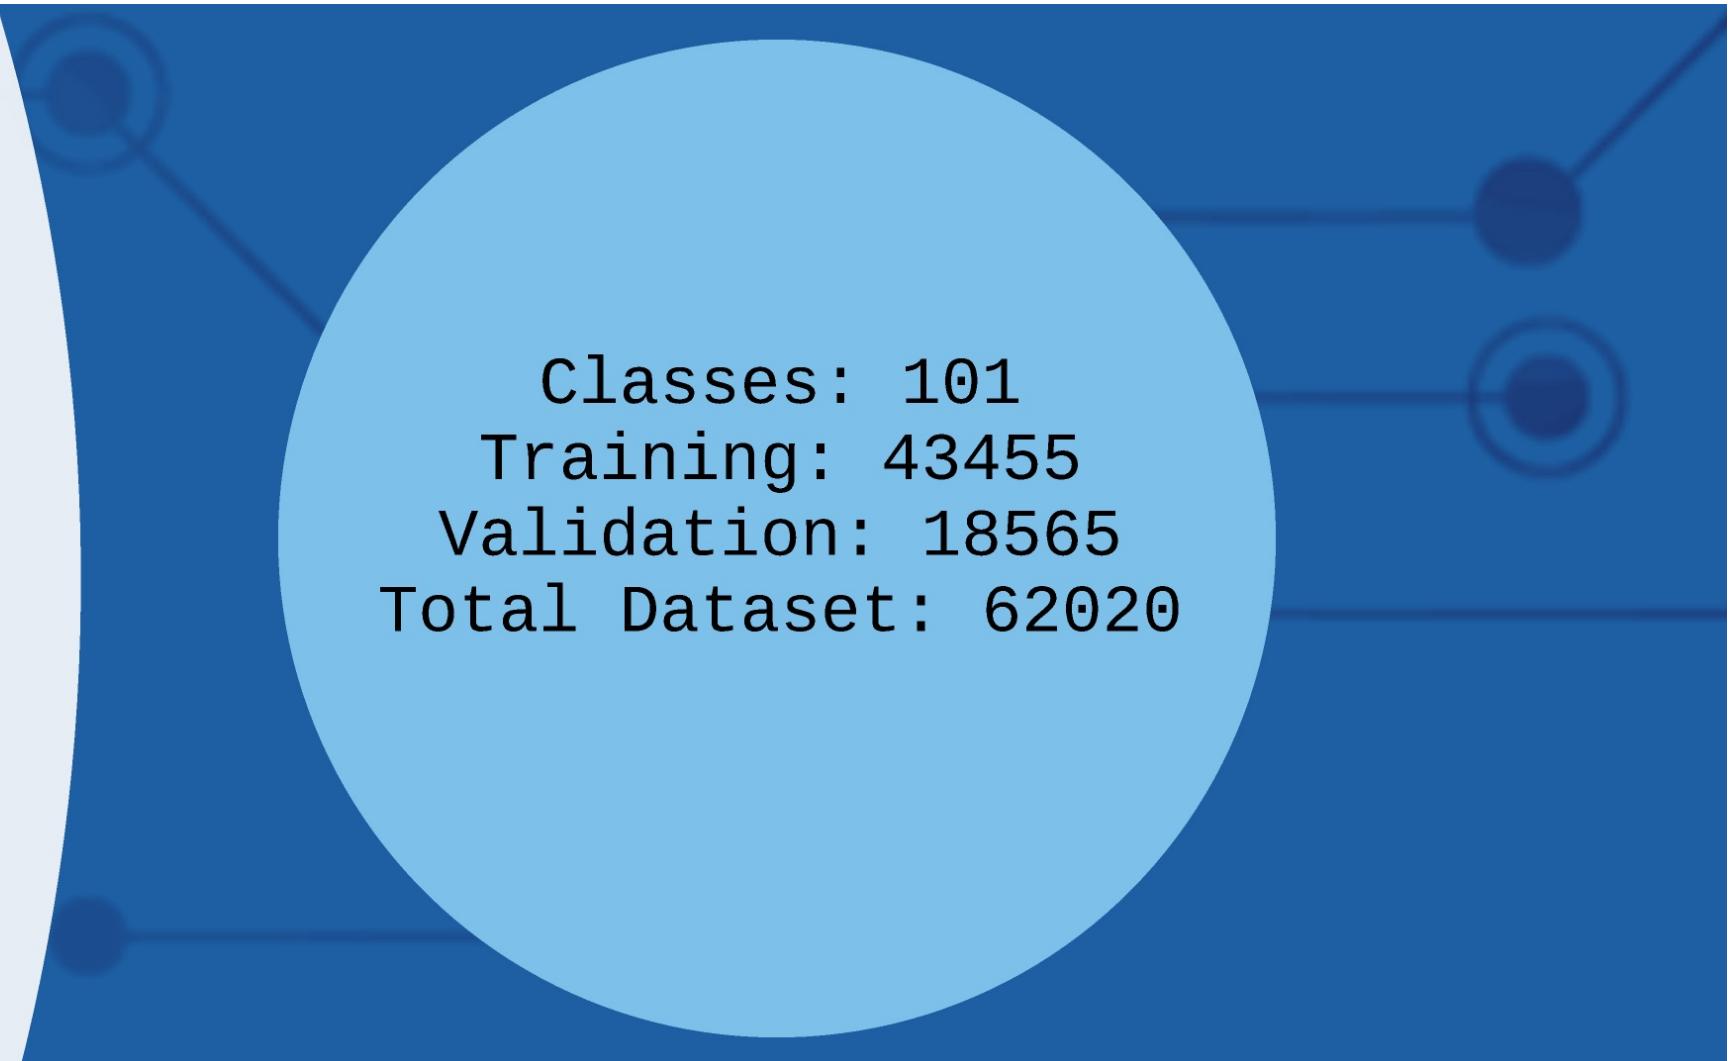


Here comes the miners:
Data Collection

Image Scraping

Amount of Data

Characteristics
of the Data



Classes: 101
Training: 43455
Validation: 18565
Total Dataset: 62020



Here comes the miners:
Data Collection

Image Scraping

Amount of Data

Characteristics
of the Data

Data Augmentation

- Rescale
- Shear Range
- Zoom Range
- Rotation range
- Horizontal Flip

**70/30 split for
Train & Test**



Here comes the miners:
Data Collection

Image Scraping

Amount of Data

Characteristics
of the Data

Introduction to Machine Learning

"The Outliers"

1.
Introduction

2.
Theoretical
Description

3.
Data Collection

4.
Models
&
Implementations

5.
Results
&
Challenges

6.
Conclusion

Implementation and fine-tuning of **Models**

GitHub: <https://tinyurl.com/TheOutliersML>

Step **01**
Example Solutions
showed
during Class

Step **02**
CNN
from Scratch

Step **03**
All
Keras Pre-trained
Models

Step **04**
Transfer Learning
Models

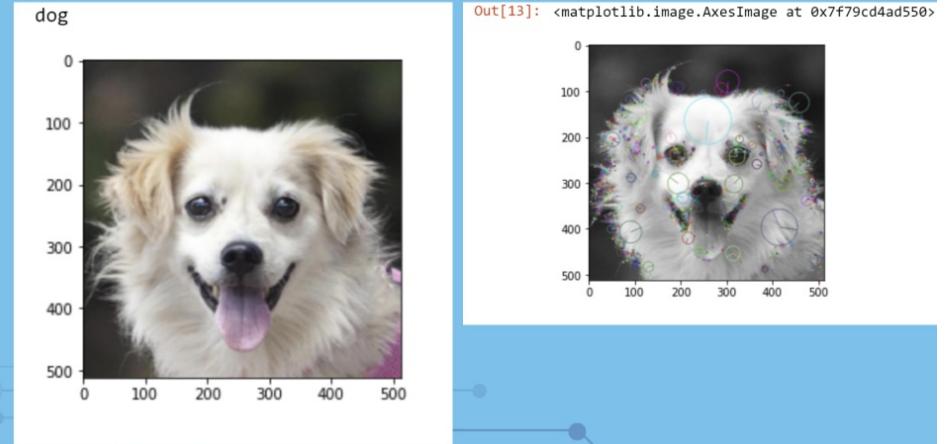
Step **05**
Additional Methods

SIFT BASED SOLUTION

- Image alignment and object detection.
- Keypoints/Interest point
- SIFT detector
- SIFT descriptor

Result on our Dataset

```
##### RESULTS #####
--> Top-1 Accuracy: 0.192
--> Top-3 Accuracy: 0.500
--> Top-10 Accuracy: 0.615
```



Implementation and fine-tuning of **Models**

GitHub: <https://tinyurl.com/TheOutliersML>

Step **01**
Example Solutions
showed
during Class

Step **02**
CNN
from Scratch

Step **03**
All
Keras Pre-trained
Models

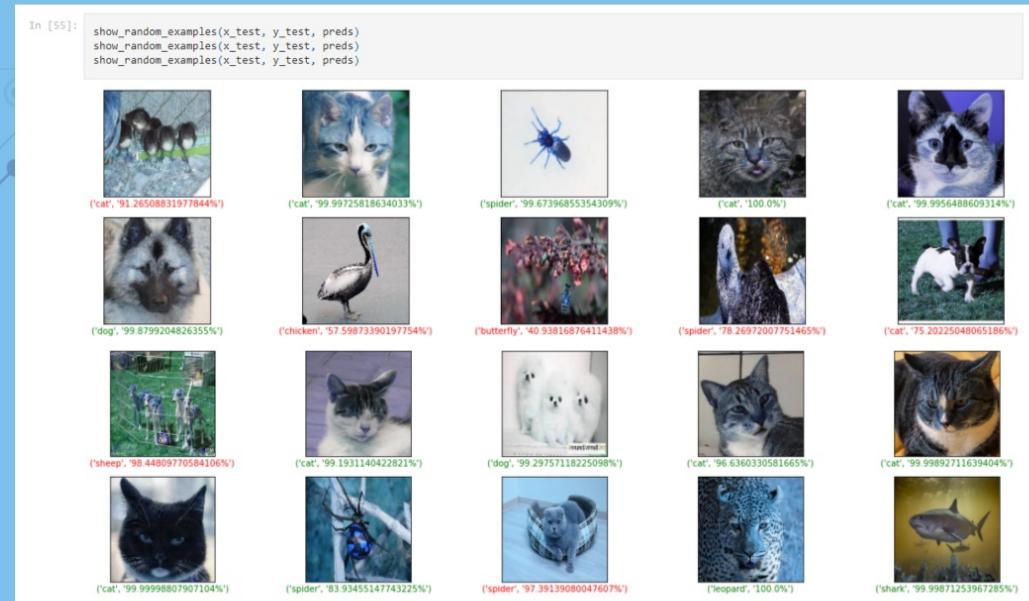
Step **04**
Transfer Learning
Models

Step **05**
Additional Methods

CNN from Scratch - 01

<https://tinyurl.com/TheOutliersCNNscratch>

- Preprocessed the data
- Converted to **Numpy Array**
- **LabelBinarizer** to scale Label data
- Train/test **70:30** split
- **Training** the model:
 - epochs=10
 - batch_size=128
 - img_size = 360

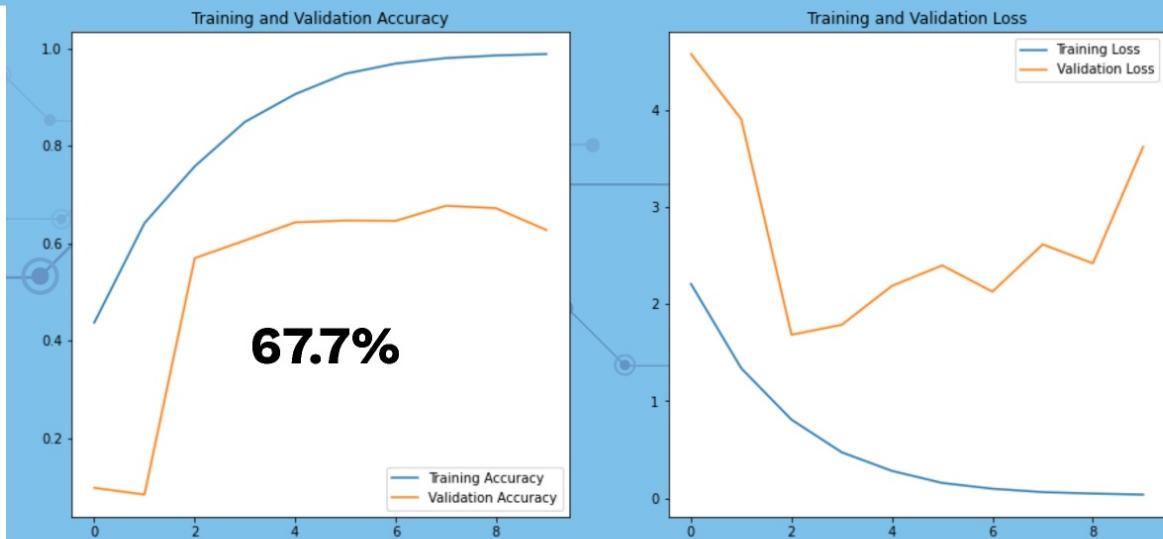


CNN from Scratch - 02

Model Summary

Model: "sequential"		
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 32)	896
batch_normalization (BatchN ormalization)	(None, 128, 128, 32)	128
conv2d_1 (Conv2D)	(None, 126, 126, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_2 (Conv2D)	(None, 63, 63, 64)	18496
batch_normalization_1 (BatchNormalization)	(None, 63, 63, 64)	256
conv2d_3 (Conv2D)	(None, 61, 61, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_4 (Conv2D)	(None, 30, 30, 128)	73856
batch_normalization_2 (BatchNormalization)	(None, 30, 30, 128)	512
conv2d_5 (Conv2D)	(None, 28, 28, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 98)	2458722
<hr/>		
Total params: 2,746,626		
Trainable params: 2,746,178		
Non-trainable params: 448		

Training and Validation



<https://tinyurl.com/TheOutliersCNNscratch>

Implementation and fine-tuning of **Models**

GitHub: <https://tinyurl.com/TheOutliersML>

Step **01**
Example Solutions
showed
during Class

Step **02**
CNN
from Scratch

Step **03**
All
Keras Pre-trained
Models

Step **04**
Transfer Learning
Models

Step **05**
Additional Methods

Keras Pre-trained Models - 01

Yes! We tried them **ALL!!!**

We are that **CRAZY!!!**

<https://keras.io/api/applications/>

The screenshot shows the Keras Applications page. At the top right, it says "33 models!" in red. Below that is a table with columns: Model, Size (MB), Top-1 Accuracy, Top-5 Accuracy, Parameters, Depth, Time (ms) per inference step (CPU), and Time (ms) per inference step (GPU). The table lists several models, including Xception, VGG16, VGG19, ResNet50, ResNet50V2, ResNet101, EfficientNetV2B1, EfficientNetV2B2, EfficientNetV2B3, EfficientNetV2S, EfficientNetV2M, and EfficientNetV2L. At the bottom of the table, there is a note about top-1 and top-5 accuracy and a note about depth.

Model	Size (MB)	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth	Time (ms) per inference step (CPU)	Time (ms) per inference step (GPU)
Xception	88	79.0%	94.5%	22.9M	81	109.4	8.1
VGG16	528	71.3%	90.1%	138.4M	16	69.5	4.2
VGG19	549	71.3%	90.0%	143.7M	19	84.8	4.4
ResNet50	98	74.9%	92.1%	25.6M	107	58.2	4.6
ResNet50V2	98	76.0%	93.0%	25.6M	103	45.6	4.4
ResNet101	171	76.4%	92.8%	44.7M	209	89.6	5.2
EfficientNetV2B1	34	79.8%	95.0%	8.2M	-	-	-
EfficientNetV2B2	42	80.5%	95.1%	10.2M	-	-	-
EfficientNetV2B3	59	82.0%	95.8%	14.5M	-	-	-
EfficientNetV2S	88	83.9%	96.7%	21.6M	-	-	-
EfficientNetV2M	220	85.3%	97.4%	54.4M	-	-	-
EfficientNetV2L	479	85.7%	97.5%	119.0M	-	-	-

The top-1 and top-5 accuracy refers to the model's performance on the ImageNet validation dataset.
Depth refers to the topological depth of the network. This includes activation layers, batch normalization layers etc.

Keras Pre-trained Models - 02

- Wrote **re-usable code templates** to run all the models
- Run with **default parameters**
- Compared **results**
 - Finalized **EfficientNetB7**
- **Fine-tuned** model:
 - batch_size = 64
 - img_size = 224
 - PCA n_components = 500
 - distance: Cosine
 - Nearest neighbor algo: KD Tree



The Keras Applications page from keras.io. The page features the Keras logo and navigation links for documentation and API reference. A table lists various pre-trained models with their specifications. The EfficientNetB7 model is highlighted.

Model	Size (MB)	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth	Time (ms) per inference step (CPU)	Time (ms) per inference step (GPU)
Xception	88	79.0%	94.5%	22.9M	81	109.4	8.1
VGG16	528	71.3%	90.1%	138.4M	16	69.5	4.2
VGG19	549	71.3%	90.0%	143.7M	19	84.8	4.4
ResNet50	98	74.9%	92.1%	25.6M	107	58.2	4.6
ResNet50V2	98	76.0%	93.0%	25.6M	103	45.6	4.4
ResNet101	171	76.4%	92.8%	44.7M	209	89.6	5.2

Model	Size (MB)	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth	Time (ms) per inference step (CPU)	Time (ms) per inference step (GPU)
EfficientNetV2B1	34	79.8%	95.0%	8.2M	-	-	-
EfficientNetV2B2	42	80.5%	95.1%	10.2M	-	-	-
EfficientNetV2B3	59	82.0%	95.8%	14.5M	-	-	-
EfficientNetV2S	88	83.9%	96.7%	21.6M	-	-	-
EfficientNetV2M	220	85.3%	97.4%	54.4M	-	-	-
EfficientNetV2L	479	85.7%	97.5%	119.0M	-	-	-

The top-1 and top-5 accuracy refers to the model's performance on the ImageNet validation dataset.
Depth refers to the topological depth of the network. This includes activation layers, batch normalization layers etc.

<https://keras.io/api/applications/>

Implementation and fine-tuning of **Models**

GitHub: <https://tinyurl.com/TheOutliersML>

Step **01**
Example Solutions
showed
during Class

Step **02**
CNN
from Scratch

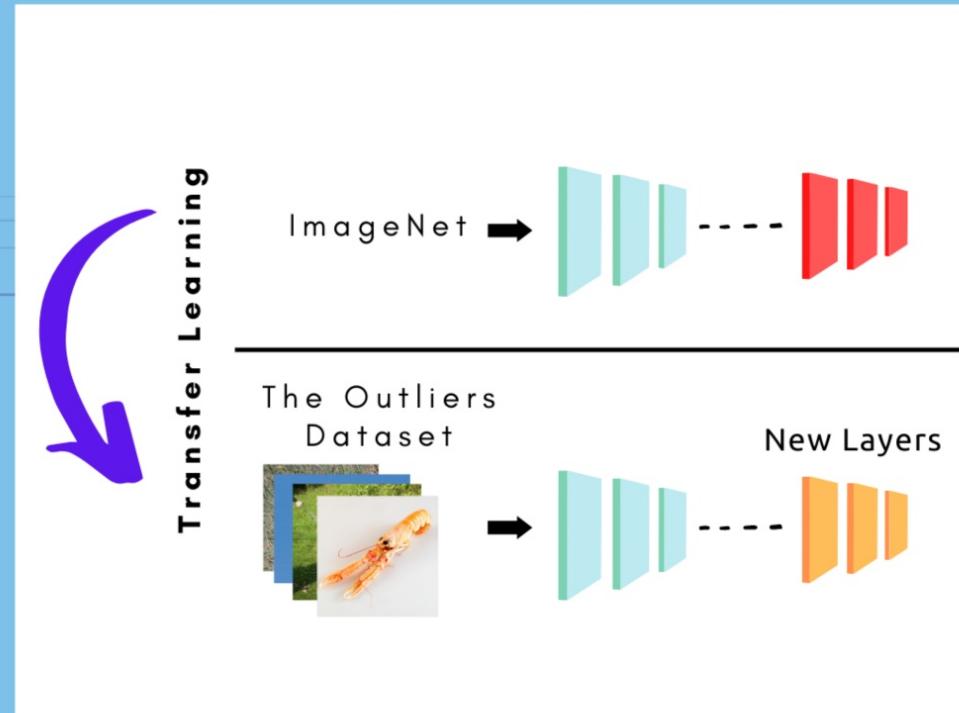
Step **03**
All
Keras Pre-trained
Models

Step **04**
Transfer Learning
Models

Step **05**
Additional Methods

Transfer Learning Models - 01

- VGG16
- Resnet50
- Resnet50v2
- MobileNet
- EfficientNetV2B3



Transfer Learning Models - 02

VGG16 TL Model: <https://tinyurl.com/TheOutliersVGG16TL>

```
# our Layers - you can add more if you want
x = Flatten()(vgg16.output)
x = Dense(1000, activation='relu')(x)
prediction = Dense(len(folders), activation='softmax')(x)

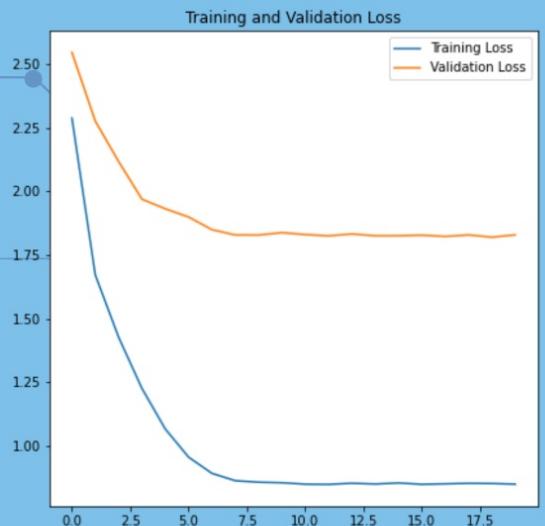
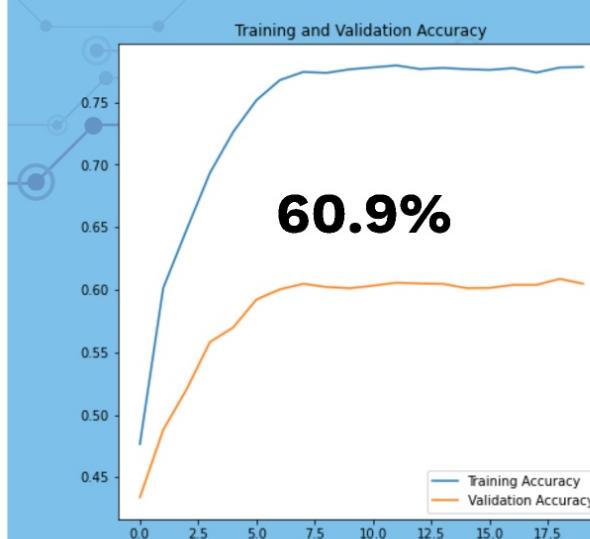
# create a model object
model = Model(inputs=vgg16.input, outputs=prediction)

# Make Last 4 Layers trainable
model.get_layer('block5_conv3').trainable = True
model.get_layer('block5_conv2').trainable = True
model.get_layer('block5_conv1').trainable = True
model.get_layer('block4_conv3').trainable = True

# view the structure of the model
model.summary()

Model: "model"
Layer (type)          Output Shape       Param #
input_1 (InputLayer)  [(None, 224, 224, 3)]   0
block1_conv1 (Conv2D)  (None, 224, 224, 64)    1792
block1_conv2 (Conv2D)  (None, 224, 224, 64)    36928
[...]
[...]
block4_conv3 (Conv2D)  (None, 28, 28, 512)    2359808
block4_pool (MaxPooling2D) (None, 14, 14, 512)  0
block5_conv1 (Conv2D)  (None, 14, 14, 512)    2359808
block5_conv2 (Conv2D)  (None, 14, 14, 512)    2359808
block5_conv3 (Conv2D)  (None, 14, 14, 512)    2359808
block5_pool (MaxPooling2D) (None, 7, 7, 512)  0
flatten (Flatten)     (None, 25088)        0
dense (Dense)         (None, 101)          2533989
[...]
Total params: 17,248,677
Trainable params: 2,533,989
Non-trainable params: 14,714,688
```

Training and Validation



Transfer Learning Models - 03

ResNet50 Model: <https://tinyurl.com/TheOutliersResNet50TL>

```
# our layers - you can add more if you want
x = Flatten()(resnet50.output)
# x = Dense(1000, activation='relu')(x)
prediction = Dense(len(folders), activation='softmax')(x)

# create a model object
model = Model(inputs=resnet50.input, outputs=prediction)

# Make last 4 layers trainable
model.get_layer('conv5_block3_3_conv').trainable = True
model.get_layer('conv5_block3_3_conv').trainable = True
model.get_layer('conv5_block3_1_conv').trainable = True
model.get_layer('conv5_block2_3_conv').trainable = True

# view the structure of the model
model.summary()

Model: "model"
Layer (type)          Output Shape       Param #  Connected to
=====
input_1 (InputLayer)   [(None, 224, 224, 3  0
)]                   []
conv1_pad (ZeroPadding2D) (None, 230, 230, 3  0
)                   ['input_1[0][0]']
conv1_conv (Conv2D)    (None, 112, 112, 64  9472
)                   ['conv1_pad[0][0]']

[...]
conv5_block3_3_conv (Conv2D)    (None, 7, 7, 2048) 1050624
['conv5_block3_2_relu[0][0]']
conv5_block3_3_bn (BatchNormal ization) (None, 7, 7, 2048) 8192
['conv5_block3_3_conv[0][0]']

conv5_block3_add (Add)     (None, 7, 7, 2048) 0
['conv5_block2_out[0][0]', 'conv5_block3_3_bn[0][0]']
conv5_block3_out (Activation) (None, 7, 7, 2048) 0
['conv5_block3_add[0][0]']
flatten (Flatten)        (None, 100352) 0
['conv5_block3_out[0][0]']
dense (Dense)            (None, 101) 10135653
['flatten[0][0]']

=====
Total params: 33,723,365
Trainable params: 10,135,653
Non-trainable params: 23,587,712
```

Training and Validation



Implementation and fine-tuning of **Models**

GitHub: <https://tinyurl.com/TheOutliersML>

Step **01**
Example Solutions
showed
during Class

Step **02**
CNN
from Scratch

Step **03**
All
Keras Pre-trained
Models

Step **04**
Transfer Learning
Models

Step **05**
Additional Methods

Additional Methods for improving performance

- Implementation of **LearningRateScheduler**

```
In [15]: from keras.callbacks import LearningRateScheduler
def lr_scheduler(epoch, lr):
    decay_rate = 0.85
    decay_step = 1
    if epoch % decay_step == 0 and epoch:
        return lr * pow(decay_rate, np.floor(epoch / decay_step))
    return lr
```

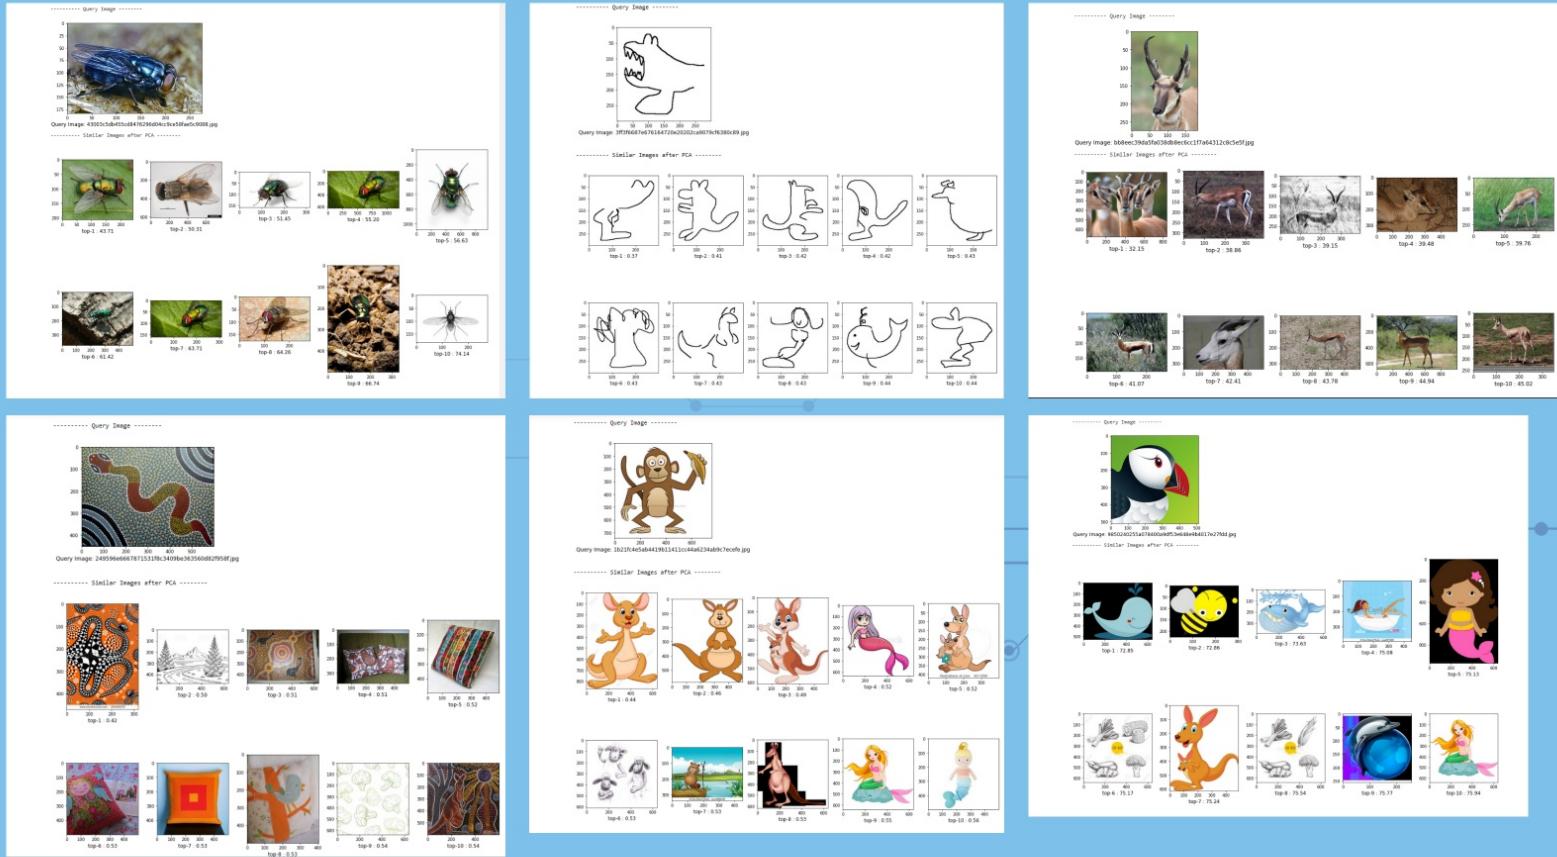
- Implementation of **PCA**

To overcome this we will apply PCA to our features and reduce the dimensions and try to find the similar features again.

```
In [27]: pca = PCA(n_components=100)
pca.fit(feature_list)
compressed_features = pca.transform(feature_list)

In [28]: neighbors_pca_features = NearestNeighbors(n_neighbors=5,
                                                algorithm='ball_tree',
                                                metric='euclidean').fit(compressed_features)
```

Some results from the Competition



Implementation and fine-tuning of **Models**

GitHub: <https://tinyurl.com/TheOutliersML>

Step **01**
Example Solutions
showed
during Class

Step **02**
CNN
from Scratch

Step **03**
All
Keras Pre-trained
Models

Step **04**
Transfer Learning
Models

Step **05**
Additional Methods

Introduction to Machine Learning

"The Outliers"

1.
Introduction

2.
Theoretical
Description

3.
Data Collection

4.
Models
&
Implementations

5.
Results
&
Challenges

6.
Conclusion



Final Result

Models we
built

Pretrained
Comparison

Best
Model

Comparison of Models we built:

Models	Top 1	Top 5	Top 10
SIFT Based Model	0.192	0.5	0.615
Image matching example	0.65	0.62	0.62
CNN from scratch	0.33	0.31	0.4
Transfer Learning on VGG16	0.4	0.59	0.597
Transfer Learning on ResNet50	0.42	0.45	0.6
Transfer Learning on MobileNet	0.5	0.62	0.69
Transfer Learning on ResNet50V2	0.49	0.5	0.58
Transfer Learning on EfficientNetV2B3	0.47	0.57	0.59





Final Result

Models we
built

Pretrained
Comparison

Best
Model

Comparison of pretrained models:

Models	Top 1	Top 5	Top 10
Transfer Learning on MobileNet	0.5	0.62	0.69
Transfer Learning on ResNet50V2	0.49	0.5	0.58
Xception	0.51	0.7	0.72
VGG16	0.53	0.56	0.62
VGG19	0.55	0.6	0.72
ResNet50V2	0.58	0.66	0.7
ResNet50	0.54	0.55	0.59
InceptionV3	0.6	0.72	0.73
MobileNet	0.62	0.63	0.7
EfficientNetB7	0.7	0.8	0.815

EfficientNetB7	Top 1	Top 5	Top 10
"Batch = 8 Img_size = 360 n_neighbors = 10 n_components = 500"	0.7	0.8	0.815
Batch = 8 Img_size = 360 n_neighbors = 15 n_components = 600	0.72	0.76	0.8
Batch = 5 Img_size = 255 n_neighbors = 15 n_components = 500	0.707	0.8	0.815
Batch = 5 Img_size = 300 n_neighbors = 10 n_components = 500	0.73	0.8	0.83



Final Result

Models we
built

Pretrained
Comparison

Best
Model

The Best Model so far:

EfficientNetB7 Pretrained Model

- **Top 1** Accuracy: **0.73**
- **Top 5** Accuracy: **0.8**
- **Top 10** Accuracy: **0.83**



Optimal Parameters:

- Batch size: 5
- image size: 300
- Distance: Cosine
- n_component: 500
- Nearest neighbor algo: KD Tree



Final Result

Models we
built

Pretrained
Comparison

Best
Model

Introduction to Machine Learning

"The Outliers"

1.
Introduction

2.
Theoretical
Description

3.
Data Collection

4.
Models
&
Implementations

5.
Results
&
Challenges

6.
Conclusion



Finally...

To Sum
Up

Are we
happy?

Possible
Improvements

It was a great experience

What did we learn?

A first hand experience in

- Data Collection
- Various ML Algorithms
- Deep Learning techniques
- Neural Networks
- Hyper-parameter tuning
- Similarity Metrics





Finally...

To Sum
Up

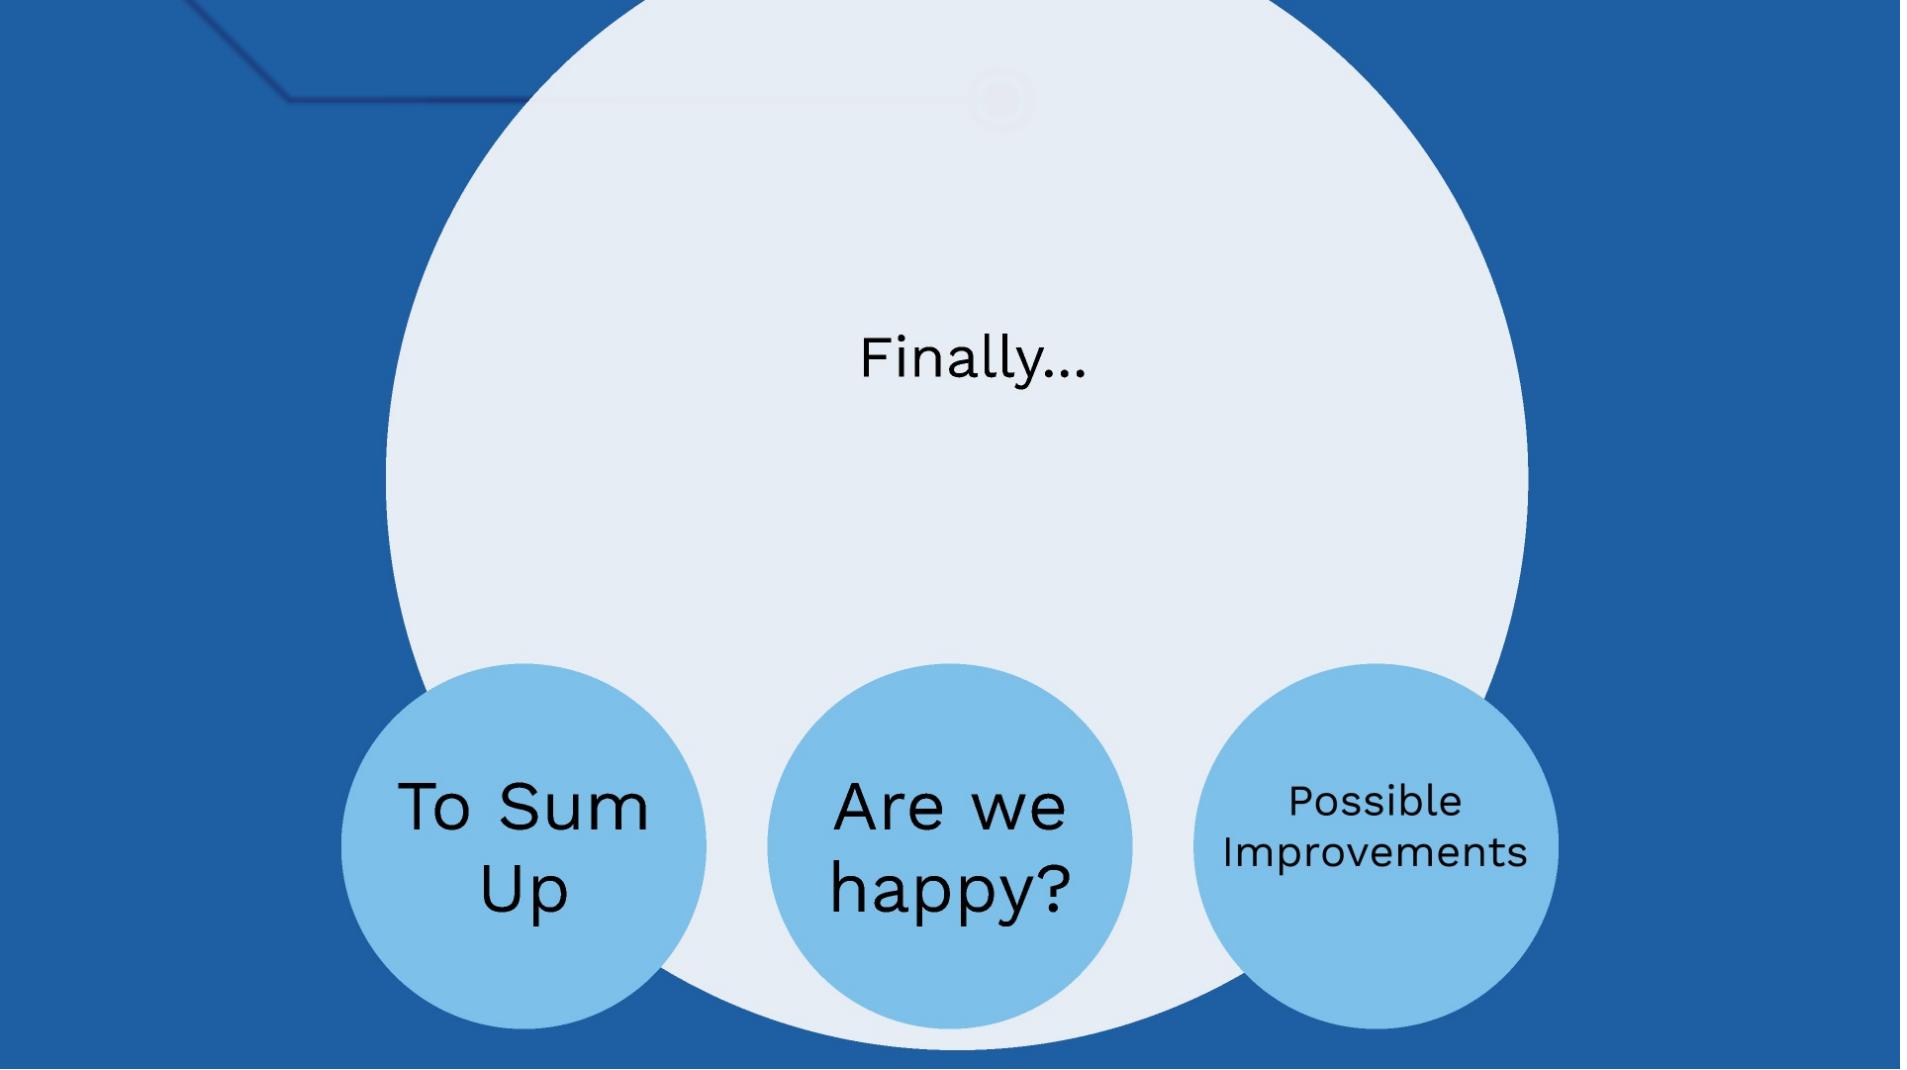
Are we
happy?

Possible
Improvements

"We can get no satisfaction"

We would have loved to have even more time to get our hands in depth, but in the end it is a project which allowed us to see how far could we get, which are our strengths and which one are our weaknesses





Finally...

To Sum
Up

Are we
happy?

Possible
Improvements

Possible improvements of our work

- Larger and more diverse Dataset
- Improved training
- Ensemble methods can be explored
- More advanced similarity methods for scalability



	Team Member	Contribution
The Outliers	Raul Fabio Riva	Content Creator/Research & Reference
		Statistical Modelling
		Fine Tuning
	Swaviman Kumar	Collaborative filtering approach
		VGG16 Transfer Model
		MobileNet Transfer Model
		Ensemble learning
	Md Ashraful Alam Hazari	ANNOY Indexing for image similarity
		Prepare Dataset
		Prepare Dataset Manually using Image downloader and Clean data
		Image Scraping
		Fine Tuning all the pre-trained models
	Shaker Mahmud Khandaker	Test the models with our own gallery and query
		Analysis and fine-tuning of the example SIFT Solution
		Implementation of CNN from Scratch
		Writing Re-usable codes to implement any models
		Implementation and fine-tuning of all Keras Pre-trained Models
		Implementation and fine-tuning of Transfer Learning Models:- VGG16- Resnet50- Resnet50v2- EfficientNetV2B3
		Implementation of LearningRateScheduler
		Implementation of PCA
		Implementation of Submission API

Thanks for the attention :)



Shaker Mahmud Khandaker

Machine Learning Engineer



Raul Fabio Riva

Research & Development,
Content Creator



Ashraful Alam Hazari

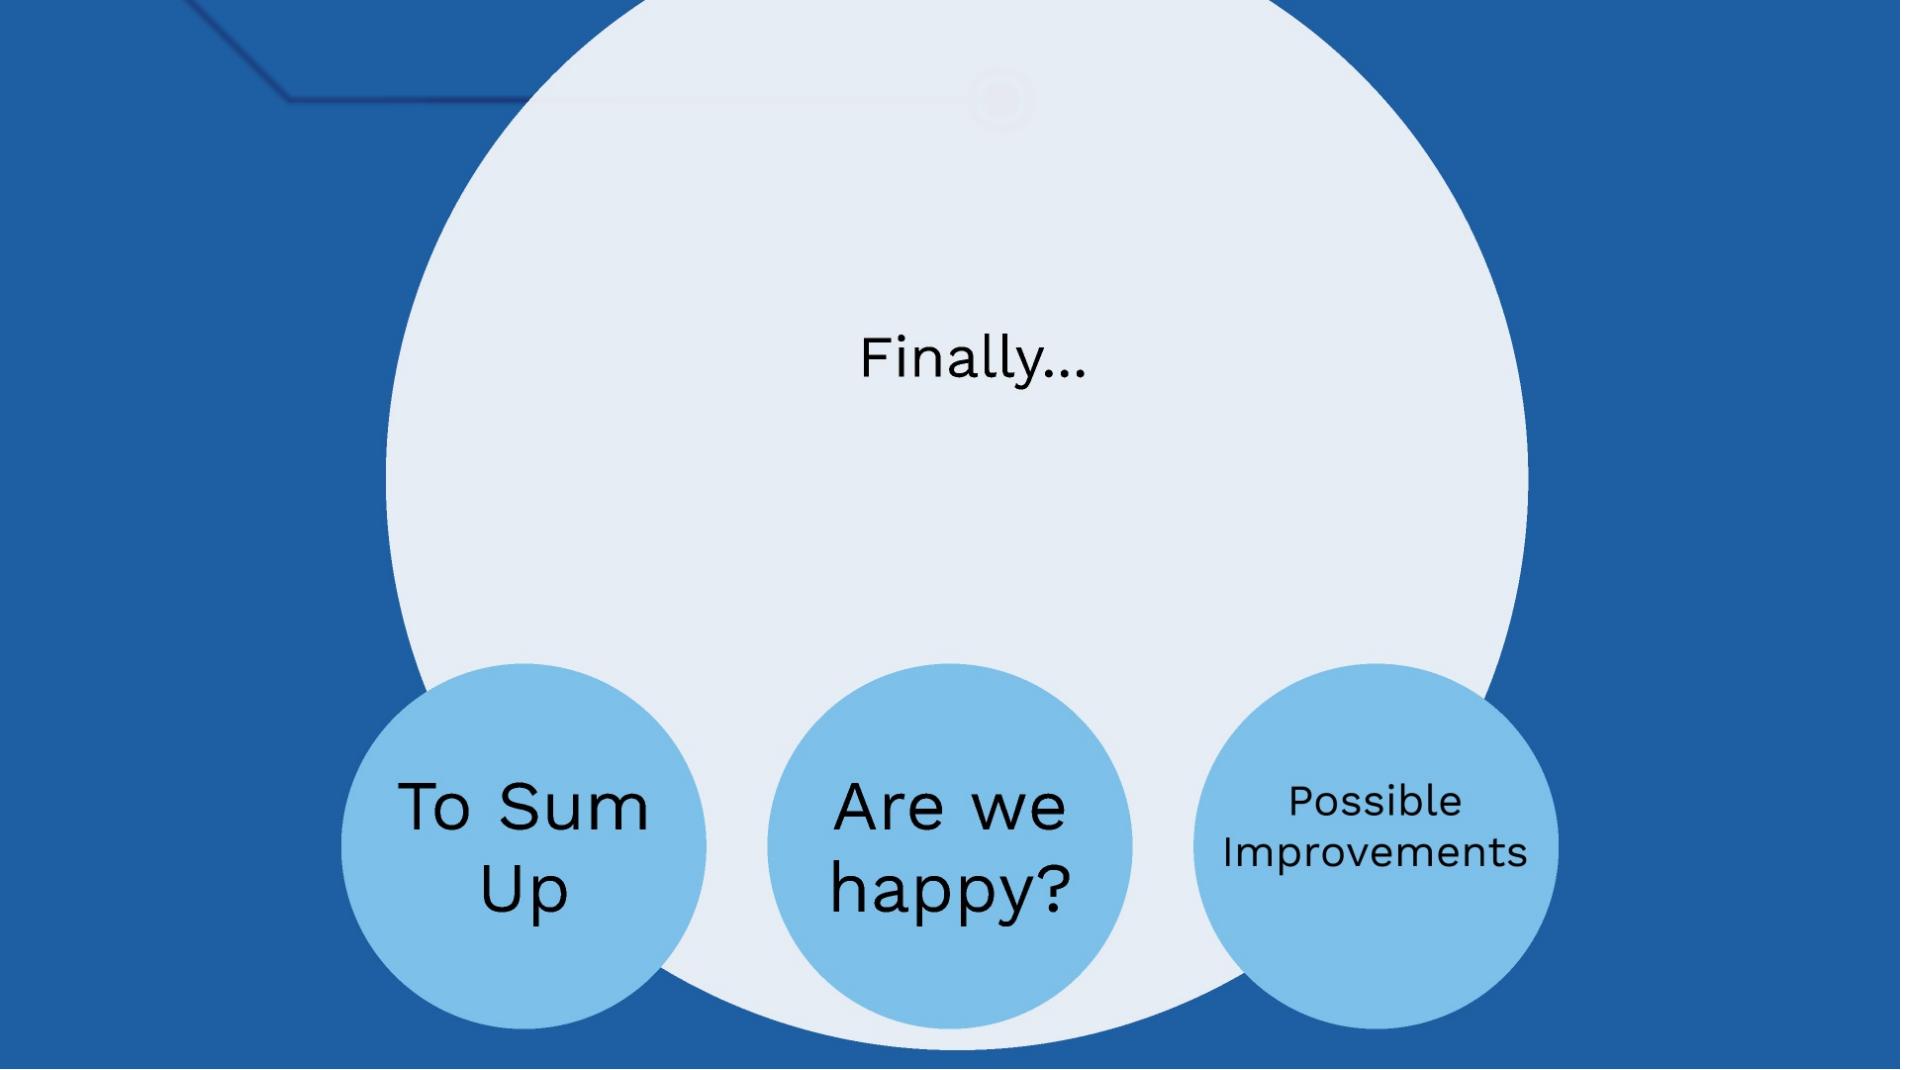
Data Miner



Swaviman Kumar

Software Engineer

Github: <https://tinyurl.com/TheOutliersML>



Finally...

To Sum
Up

Are we
happy?

Possible
Improvements

Introduction to Machine Learning

"The Outliers"

1.
Introduction

2.
Theoretical
Description

3.
Data Collection

4.
Models
&
Implementations

5.
Results
&
Challenges

6.
Conclusion